

R and Python clients for GBIF species occurrence data

Scott Chamberlain^{*,a}, Carl Boettiger^b

^a*University of California, Berkeley, CA, USA*

^b*University of California, Berkeley, CA, USA*

Abstract

1. xxx
2. xxx
3. xxx
4. xxxx

Introduction

Perhaps the most fundamental element in many fields of ecology is the individual organism. The number of individuals of each species in a given location forms the basis for many sub-fields of ecology and evolution. Some research questions necessitate collecting new data, while others can easily take advantage of existing data. In fact, some ecology fields are built largely on existing data, e.g., macro-ecology (Brown, 1995; Beck et al., 2012).

Data on individuals, including which species, and where they're found, can be used for a large number of research questions. Biodiversity records have been used for a suite of other use cases: validating habitat suitability models with real occurrence data (Ficetola et al., 2014); ancestral range reconstruction (Ferretti et al., 2015; María Mendoza et al., 2015); development of invasive species watch lists (Faulkner et al., 2014); evaluating risk of invasive species spread (Febbraro et al., 2013); and effects of climate change on future biodiversity (Brown et al., 2015).

In addition to wide utility, this data is important for conservation. Biodiversity loss is one of the greatest challenges of our time (Pimm et al., 2014), and some have called this the sixth great mass extinction (Ceballos et al., 2015). Given this challenge there is a great need for data on specimen records, whether collected from live sightings in the field or specimens in museums.

*Corresponding author

Email addresses: `scott(at)ropensci.org` (Scott Chamberlain), `carl(at)ropensci.org` (Carl Boettiger)

26 Global Biodiversity Information Facility

27 There are many online services that collect and maintain specimen records. However, Global Biodiversity
28 Information Facility (hereafter, GBIF, <http://www.gbif.org>) is the largest collection of biodiversity
29 records globally, currently with 643 million records, 1.6 million taxa, 15,450 datasets from 780 publishers
30 (as of 2016-02-09). Many large biodiversity warehouses such as iNaturalist (<http://www.inaturalist.org>),
31 VertNet (<http://vertnet.org>), and USGS's Biodiversity Information Serving Our Nation (BISON;
32 <http://bison.usgs.ornl.gov>) all feed into GBIF.

33 The most important data organizational level in GBIF is the individual record. In R using `rgbif`
34 you'll recognize this as a row in the data.frame that is returned from `occ_search()` or `occ_data()`.
35 Going upstream, each record is part of a dataset, where each dataset is submitted by an organization,
36 organizations are organized into nodes, datasets are published through institutions (which may be
37 hosted at another organization), and a network is a group of datasets (managed by GBIF).

38 Each record has some taxonomic name associated with it, which itself is linked to a lot of other
39 taxonomic data. GBIF maintains their own taxonomic data.

40 All this organization matters because you can navigate through GBIF data through all these organiza-
41 tional levels.

42 The clients

43 *The rgbif package*

44 Herein, we describe the `rgbif` software package (Chamberlain et al.) for working with GBIF data in the
45 R programming environment (R Core Team, 2014). R is a widely used language in academia, as well as
46 non-profit and private sectors. Importantly, R makes it easy to execute all steps of the research process,
47 including data management, data manipulation and cleaning, statistics, and visualization. Thus, an R
48 client for getting GBIF data is a powerful tool to facilitate reproducible research.

49 The `rgbif` package is nearly completely written in R (a small Javascript library is included for
50 reading well known text (Herring, 2011)), uses an [MIT license](#) to maximize use everywhere. `rgbif` is
51 developed publicly on GitHub at <https://github.com/ropensci/rgbif>, where development versions of
52 the package can be installed, and bugs and feature requests reported. Stable versions of `rgbif` can be
53 installed from [CRAN](#), the distribution network for R packages. `rgbif` is part of the rOpenSci project
54 (<http://ropensci.org>), a developer network making R software to facilitate reproducible research.

55 *The pygbif library*

56 `pygbif` (Chamberlain) is a Python library for working with GBIF data in the Python programming
57 environment (???). Python is a general purpose programming language used widely in all sectors,
58 and for all parts of software development including server and client side use cases. Python is used
59 exclusively in some scientific disciplines (e.g., astronomy), and has partial usage in other disciplines. A
60 Python client for GBIF data is an important tool given the even wider usage of Python than R, though
61 maybe slightly less than R for ecology/biology.

```
pip install pygbif
```

```
import pygbif
```

62 The `pygbif` library is less mature and complete than the R package. It also uses an [MIT license](#) to
63 maximize use everywhere. `pygbif` is developed publicly on GitHub at <https://github.com/sckott/pygbif>,
64 where development versions of the package can be installed, and bugs and feature requests reported.
65 Stable versions of `pygbif` can be installed from [pypi](#), the distribution network for Python libraries.

66 *Library interfaces*

67 `rgbif` and `pygbif` are designed following the [GBIF Application Programming Interface](#), or API. The
68 GBIF API has four major components: registry, taxonomic names, occurrences, and maps. We also
69 include functions to interface with the OAI-PMH GBIF service; only dataset information is available
70 vis this service, however. We ignore maps in both libraries as it is concerned with generating maps
71 primarily for web applications. Both libraries have a suite of functions dealing with each of registry,
72 taxonomic names, and occurrences - we'll go through each in turn describing design and example usage.

73 *GBIF headers*

74 With each request `rgbif` and `pygbif` make to GBIF's API, we send request headers that tell GBIF what
75 client the request is coming from, including what version of the library. This helps GBIF know what
76 proportion of requests are coming from which client, and therefore from R vs. Python; this information
77 is helpful for GBIF in thinking about how people are using GBIF data.

78 *Registry*

79 The GBIF registry API services are spread across five sets of functions via the main GBIF API:

- 80 • Datasets
- 81 • Installations
- 82 • Networks
- 83 • Nodes
- 84 • Organizations

85 Dataset information in general is available via the OAI-PMH service, functions in `rgbif` prefixed with
86 `gbif_oai_`, but not available in `pygbif` yet.

87 Datasets are owned by organizations. Organizations are endorsed by nodes to share datasets with GBIF.
88 Datasets are published through institutions, which may be hosted at another organization. A network
89 is a group of datasets (managed by GBIF). Datasets are the units that matter the most with respect
90 to registry information, while installations, networks, nodes, and organizations are simply higher level
91 organizational structure.

92 *Datasets*

93 Dataset functions include search, dataset metadata retrieval, and dataset metrics. Searching for datasets
94 is an important part of the discovery process. One can search for datasets on the GBIF web portal.
95 However, programmatic searching using this package is much more powerful. Identifying datasets
96 appropriate for a research question is helpful as you can get metadata for each dataset, and track down
97 dataset specific problems, if any.

98 The `dataset_search()` function is one way to search for datasets. Here, we search for the query term
99 “oregon”, which finds any datasets that have terms matching that term.

```
res <- dataset_search(query = "oregon")
res$data$datasetTitle[1:10]
#> [1] "UCDavis - Western USA - Monarch Butterflies - 1892-2005"
#> [2] "A geographic distribution database of Mononychellus mites (Acari: Tetranychidae) on cas
#> [3] "Plantas Acuáticas de la Orinoquía Colombiana"
```

```
#> [4] "USBombus, contemporary survey data of North American bumble bees (Hymenoptera, Apidae,
#> [5] "Wool carder bees of the genus Anthidium in the Western Hemisphere"
#> [6] "CM Birds Collection"
#> [7] "SDNHM Birds Collection"
#> [8] "University of British Columbia Herbarium (UBC) - Bryophytes Collection"
#> [9] "University of British Columbia Herbarium (UBC) - Vascular Plant Collection"
#> [10] "Bryophyte Collection - University of Washington Herbarium (WTU)"
```

100 See also `datasets()` and `dataset_suggest()` in `rgbif` for searching for datasets.

101 In Python, we can similarly search for datasets. Here, search for datasets of type `OCCURRENCE`:

```
from pygbif import registry
registry.datasets(type="OCCURRENCE")
```

102 *Dataset metrics.* Dataset metrics are another useful way of figuring out what datasets you may want to
 103 use. One drawback is that these metrics data are only available for datasets of type *checklist*, but there
 104 are quite a lot of them (2737).

105 Here, we search for dataset metrics for a single dataset, with uuid `ec93a739-1681-4b04-b62f-3a687127a17f`,
 106 a checklist of the ants (Hymenoptera: Formicidae) of the World.

```
res <- dataset_metrics(uuid='ec93a739-1681-4b04-b62f-3a687127a17f')
data.frame(rank = names(res$countByRank),
           count = unname(unlist(res$countByRank)))
```

rank	count
SPECIES	13710
SUBSPECIES	3234
GENUS	726
TRIBE	53
SUBFAMILY	20
FAMILY	2
KINGDOM	1

rank	count
PHYLUM	1
CLASS	1
ORDER	1

107 And in Python, get metrics for the same dataset as above:

```
from pygbif import registry
registry.dataset_metrics(uuid='ec93a739-1681-4b04-b62f-3a687127a17f')
```

108 *Networks, nodes, and installations*

109 Networks, nodes and installations are at a higher level of organization above datasets, but can be
 110 useful if you want to explore data from given organizations. Here, in R we search for the first 10 GBIF
 111 networks, returning just the title field.

```
networks(limit = 10)$data$title
#> [1] "GBIF Backbone Sources"
#> [2] "Canadensys"
#> [3] "Southwest Collections of Arthropods Network (SCAN)"
#> [4] "VertNet"
#> [5] "Dryad"
#> [6] "GBIF Network"
#> [7] "The Knowledge Network for Biocomplexity (KNB) "
#> [8] "Online Zoological Collections of Australian Museums (OZCAM)"
#> [9] "Catalogue of Life"
#> [10] "Ocean Biogeographic Information System (OBIS)"
```

112 And in Python:

```
from pygbif import registry
registry.networks(limit = 10)
```

113 *Taxonomic names*

114 The GBIF taxonomic names API services are spread across five functions in `rgbif`:

- 115 • Search GBIF name backbone - `name_backbone()`
- 116 • Search across all checklists - `name_lookup()`
- 117 • Quick name lookup - `name_suggest()`
- 118 • Name usage of a name according to a checklist - `name_usage()`
- 119 • GBIF name parser - `parsenames()`

120 `pygbif` only has `name_backbone()` and `name_suggest()` at this time.

121 The goal of these name functions is often to settle on a taxonomic name known to GBIF's database.
122 This serves two purposes: 1) when referring to a taxonomic name, you can point to a URI on the
123 Internet, and 2) you can search for metadata on a taxon, and occurrences of that taxon in GBIF.

124 Taxonomic names are particularly tricky. Many different organizations have their own unique codes for
125 the same taxonomic names, and some taxonomic groups have preferred sources for the definitive names
126 for that group. That's why it's best to determine what name GBIF uses, and its associated identifier,
127 for the taxon of interest instead of simply searching for occurrences with a taxonomic name.

128 When searching for occurrences (see below) you can search by taxonomic name (and other filters, e.g.,
129 taxonomic rank), but you're probably better off figuring out the taxonomic key in the GBIF backbone
130 taxonomy, and using that to search for occurrences. The `taxonkey` parameter in the GBIF occurrences
131 API expects a GBIF backbone taxon key.

132 *GBIF Backbone*

133 The GBIF backbone taxonomy is used in GBIF to have a consistent way to refer to taxonomic
134 names throughout their services. The backbone has 4410899 unique names and 2497114 species
135 names. The backbone taxonomy is also a dataset with key `d7dddbf4-2cf0-4f39-9b2a-bb099caae36c`
136 (<http://www.gbif.org/dataset/d7dddbf4-2cf0-4f39-9b2a-bb099caae36c>).

137 We can search the backbone taxonomy with the function `name_backbone()` in both R and Python
138 clients. Here, we're searching for the name *Poa*, restricting to genera, and the family *Poaceae*, in R

```

res <- name_backbone(name='Poa', rank='genus', family='Poaceae')
res[c('usageKey', 'kingdom')]
#> $usageKey
#> [1] 2704173
#>
#> $kingdom
#> [1] "Plantae"

```

139 and in Python

```

from pygbif import species
res = species.name_backbone(name='Poa', rank='genus', family='Poaceae')
[ res[x] for x in ['usageKey', 'kingdom'] ]

```

140 *Name searching*

141 One of the quickest ways to search for names is using `name_suggest()`, which does a very quick search
 142 and returns minimal data. Here, we're searching for the query term *Pum*, and we get back many names:

```
name_suggest(q='Pum', limit = 6)
```

key	canonicalName	rank
4848380	Pumiliornis	GENUS
2235094	Pumilibranchipus	GENUS
1795637	Pumora	GENUS
4581745	Pumicia	GENUS
4598790	Pumilicopta	GENUS
1593095	Pumilomyia	GENUS

143 The same in Python


```
from pygbif import species
species.name_suggest(q='Pum', limit = 6)
```

With these results, you can then proceed to search for occurrences with the taxon key(s), or drill down further with other name searching functions to get the exact taxon of interest.

Occurrences

GBIF provides two ways to get occurrence data: through the `/occurrence/search` route (see `occ_search` in `rgbif`, or `occurrences.search` in `pygbif`), or via the `/occurrence/download` route (many functions, see below). `occ_search()`/`occurrences.search` is the main function for the search route, and is more appropriate for smaller data, while the download functions are more appropriate for larger data requests.

GBIF imposes for any given search a limit of 200,000 records in the search service, after which point you can't download any more records for that search. However, you can download more records for different searches.

We think the search service is still quite useful for many people even given the 200,000 limit. For those that need more data, we have created a similar interface in the download functions, that should be easy to use. Users should take note that using the download service has a few extra steps to get data into R, but is straight-forward.

The download service, like the occurrence search service, is rate-limited. That is, you can only have one to three downloads running simultaneously for your user credentials. However, simply check when a download job is complete, then you should be able to start a new download request.

Download API

The download API syntax is similar to the occurrence search API in that the same parameters are used, but the way in which the query is defined is different. For example, in the download API you can do greater than searches (i.e., `latitude > 50`), whereas you cannot do that in the occurrence search API. Thus, unfortunately, we couldn't make the query interface exactly the same for both search and download functions.

Using the download service can consist of as few as three steps: 1) Request data via a search; 2) Download data; 3) Import data into R.

170 Request data download given a query. Here, we search for the taxon key 3119195, which is the key for
171 *Helianthus annuus* (<http://www.gbif.org/species/3119195>).

```
occ_download('taxonKey = 3119195')  
#> <<gbif download>>  
#> Username: xxxx  
#> E-mail: xxxx  
#> Download key: 0000840-150615163101818
```

172 You can check on when the download is ready using the functions `occ_download_list()` and
173 `occ_download_meta()`. When it's ready use `occ_download_get()` to download the dataset to your
174 computer.

```
(res <- occ_download_get("0000840-150615163101818", overwrite = TRUE))  
#> <<gbif downloaded get>>  
#> Path: ./0000840-150615163101818.zip  
#> File size: 3.19 MB
```

175 What's printed out above is a very brief summary of what was downloaded, the path to the file, and its
176 size (in human readable form).

177 Next, read the data in to R using the function `occ_download_import()`.

```
library("dplyr")  
dat <- occ_download_import(res)  
dat %>%  
  select(gbifID, decimalLatitude, decimalLongitude)  
#>      gbifID abstract accessRights accrualMethod accrualPeriodicity accrualPolicy alternative  
#> 1  725767384      NA              NA              NA              NA              NA  
#> 2  725767447      NA              NA              NA              NA              NA  
#> 3  725767450      NA              NA              NA              NA              NA  
#> 4  725767513      NA              NA              NA              NA              NA  
#> 5  725767546      NA              NA              NA              NA              NA  
#> 6  725767579      NA              NA              NA              NA              NA
```

```
#> 7 725767609      NA      NA      NA      NA      NA
#> 8 725767645      NA      NA      NA      NA      NA
#> 9 725767678      NA      NA      NA      NA      NA
#> 10 725767681     NA      NA      NA      NA      NA
#> ..      ...      ...      ...      ...      ...
#> Variables not shown: available (lgl), bibliographicCitation (chr), conformsTo (lgl), contribu
#> coverage (lgl), created (chr), creator (lgl), date (lgl), dateAccepted (lgl), dateCopyri
#> (lgl), dateSubmitted (lgl), description (lgl), educationLevel (lgl), extent (lgl), forma
#> hasFormat (lgl), hasPart (lgl), hasVersion (lgl), identifier (chr), instructionalMethod
```

178 In Python

```
from pygbif import occurrences as occ
occ.download('taxonKey = 3119195')
(res = occ.download_get("0000840-150615163101818", overwrite = True))
```

179 We don't have pygbif functionality at the moment for importing data, but it's coming soon.

180 *Downloaded data format.* The downloaded dataset from GBIF is a Darwin Core Archive (DwC-A), an
 181 internationally recognized biodiversity informatics standard (<http://rs.tdwg.org/dwc/>). The DwC-A
 182 downloaded is a compressed folder with a number of files, including metadata, citations for each of the
 183 datasets included in the download, and the data itself, in separate files for each dataset as well as one
 184 single .txt file. In `rgbif::occ_download_import()`, we simply fetch data from the .txt file. If you
 185 want to dig into the metadata, citations, etc., it is easily accessible from the folder on your computer.

186 *Search API*

187 The search API follows the GBIF API and is broken down into the following functions:

- 188 • Get a single numeric count of occurrences - `rgbif: occ_count()` / `pygbif: occurrences.count`
- 189 • Search for occurrences - `rgbif: occ_search()` / `pygbif: occurrences.count`
- 190 • A simplified and optimized version of `rgbif: occ_search()` or `occ_data()` / `pygbif:`
 191 `occurrences.count`
- 192 • Get occurrences by occurrence identifier - `rgbif: occ_get()` / `pygbif: occurrences.count`
- 193 • Get occurrence metadata - `rgbif: occ_metadata()` / `pygbif: occurrences.count`

194 *Search for occurrences.* **R**

195 The main search work-horse is `occ_search()`. This function allows very flexible search definitions. In
196 addition, this function does paging internally, making it such that the user does not have worry about
197 the 300 records per request limit - but of course we can't go over the 200,000 maximum limit.

198 The output of `occ_search()` presents a compact **data.frame** so that no matter how large the
199 **data.frame**, the output is easily assessed because only a few of the records (rows) are shown, only a few
200 columns are shown (with others shown in name only), and metadata is shown on top of the **data.frame**
201 to indicate data found and returned, media records found, unique taxonomic hierarchies returned, and
202 the query executed.

203 The output of these examples, except one, aren't shown, but all run correctly.

204 Search by species name, using `name_backbone()` first to get key

```
(key <- name_suggest(q = 'Helianthus annuus', rank = 'species')$key[1])
#> [1] 3119195
occ_search(taxonKey = key, limit = 2)
#> Records found [21538]
#> Records returned [2]
#> No. unique hierarchies [1]
#> No. media records [2]
#> Args [taxonKey=3119195, limit=2, offset=0, fields=all]
#> First 10 rows of data
#>
#>           name           key decimalLatitude decimalLongitude
#> 1 Helianthus annuus 1249279611          34.04810         -117.79884
#> 2 Helianthus annuus 1249286909          32.58747          -97.10081
#> Variables not shown: issues (chr), datasetKey (chr), publishingOrgKey
#> (chr), publishingCountry (chr), protocol (chr), lastCrawled (chr),
#> lastParsed (chr), extensions (chr), basisOfRecord (chr), taxonKey
#> (int), kingdomKey (int), phylumKey (int), classKey (int), orderKey
#> (int), familyKey (int), genusKey (int), speciesKey (int),
#> scientificName (chr), kingdom (chr), phylum (chr), order (chr),
```

```
#> family (chr), genus (chr), species (chr), genericName (chr),
#> specificEpithet (chr), taxonRank (chr), dateIdentified (chr), year
#> (int), month (int), day (int), eventDate (chr), modified (chr),
#> lastInterpreted (chr), references (chr), identifiers (chr), facts
#> (chr), relations (chr), geodeticDatum (chr), class (chr), countryCode
#> (chr), country (chr), rightsHolder (chr), identifier (chr),
#> verbatimEventDate (chr), datasetName (chr), gbifID (chr),
#> verbatimLocality (chr), collectionCode (chr), occurrenceID (chr),
#> taxonID (chr), license (chr), recordedBy (chr), catalogNumber (chr),
#> http...unknown.org.occurrenceDetails (chr), institutionCode (chr),
#> rights (chr), eventTime (chr), identificationID (chr),
#> coordinateAccuracy (dbl), coordinateAccuracyInMeters (dbl),
#> occurrenceRemarks (chr)
```

205 Instead of getting a taxon key first, you can search for a name directly

```
occ_search(scientificName = 'Ursus americanus')
```

206 Search for many species

```
splist <- c('Cyanocitta stelleri', 'Junco hyemalis', 'Aix sponsa')
keys <- sapply(splist, function(x) name_suggest(x)$key[1], USE.NAMES = FALSE)
occ_search(taxonKey = keys, limit = 5, return = 'data')
```

207 Spatial search, based on well known text format (Herring, 2011), or a bounding box set of four coordinates

```
# well known text
occ_search(geometry = 'POLYGON((30.1 10.1, 10 20, 20 40, 40 40, 30.1 10.1))', limit = 20)
# bounding box
occ_search(geometry = c(-125.0,38.4,-121.8,40.9), limit = 20)
```

208 Get only occurrences with lat/long data

```
occ_search(hasCoordinate = TRUE, limit = 5)
```

209 Get only those occurrences with spatial issues. Spatial issues are a set of issues that are returned in
210 the `issues` field. They each indicate something different about that record. For example, the issue
211 `COUNTRY_COORDINATE_MISMATCH` indicates that the interpreted occurrence coordinates fall outside of
212 the indicated country. You can see how that might be useful when it comes to cleaning your data prior
213 to analysis/visualization.

```
occ_search(hasGeospatialIssue = TRUE, limit = 5)
```

214 Python

215 The equivalent occurrence search workhorse in `pygibf` is `occurrences.search`.

216 Search by species name, using `name_backbone()` first to get key

```
from pygibf import species
from pygibf import occurrences as occ
key = species.name_suggest(q = 'Helianthus annuus', rank = 'species')['data'][0]['key']
occ.search(taxonKey = key, limit = 2)
```

217 Instead of getting a taxon key first, you can search for a name directly

```
occ.search(scientificName = 'Ursus americanus')
```

218 Search for many species

```
from pygibf import species
from pygibf import occurrences as occ
splist = ['Cyanocitta stelleri', 'Junco hyemalis', 'Aix sponsa']
keys = [ species.name_suggest(x)['data'][0]['key'] for x in splist ]
occ.search(taxonKey = keys, limit = 5)
```

219 Spatial search, based on well known text format (Herring, 2011), or a bounding box set of four coordinates

```

from pygbif import occurrences as occ
# well known text
occ.search(geometry = 'POLYGON((30.1 10.1, 10 20, 20 40, 40 40, 30.1 10.1))', limit = 20)
# bounding box
occ.search(geometry = '-125.0,38.4,-121.8,40.9', limit = 20)

```

220 Get only occurrences with lat/long data

```

from pygbif import occurrences as occ
occ.search(hasCoordinate = True, limit = 5)

```

221 Get only those occurrences with spatial issues. Spatial issues are a set of issues that are returned in
 222 the `issues` field. They each indicate something different about that record. For example, the issue
 223 `COUNTRY_COORDINATE_MISMATCH` indicates that the interpreted occurrence coordinates fall outside of
 224 the indicated country. You can see how that might be useful when it comes to cleaning your data prior
 225 to analysis/visualization.

```

from pygbif import occurrences as occ
occ.search(hasGeospatialIssue = True, limit = 5)

```

226 *Data cleaning.* GBIF provides optional data issues with each occurrence record. These issues fall into
 227 many different pre-defined classes, covering issues with taxonomic names, geographic data, and more
 228 (see `occ_issues_lookup()` to find out more information on GBIF issues; and the same data on [GBIF's](#)
 229 [development site](#)).

230 `occ_issues()` provides a way to easily filter data downloaded via `occ_search()` based on GBIF issues.

```

out <- occ_search(issue = 'DEPTH_UNLIKELY', limit = 500)
NROW(out)
#> [1] 4
out %>% occ_issues(-cudc) %>% .$data %>% NROW
#> [1] 4

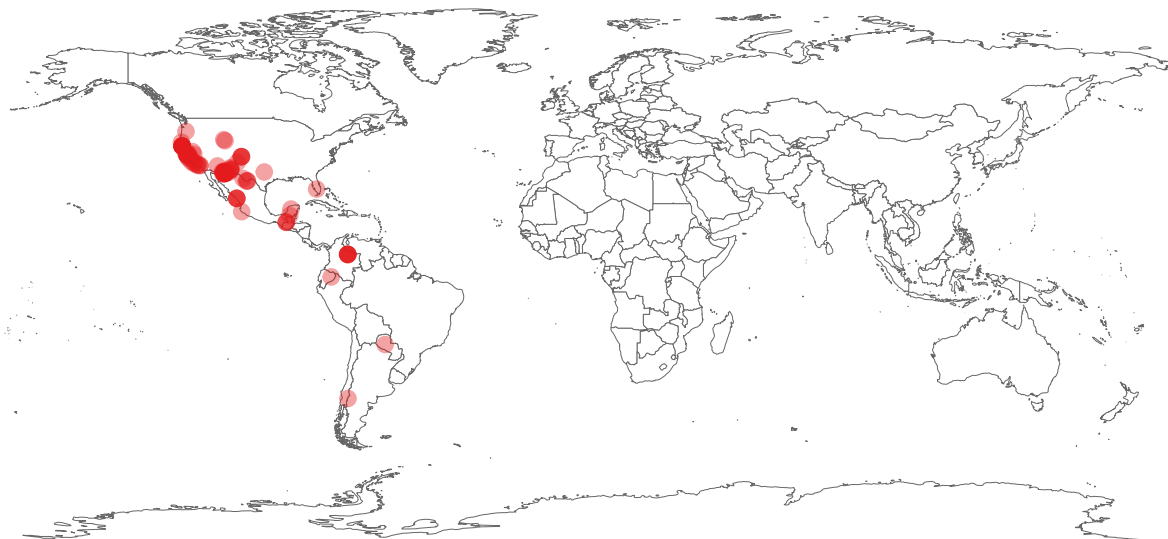
```

231 There's no equivalent interface in `pygbif`.

232 Mapping

233 An obvious downstream use case for species occurrence data is to map the data. `rgbif` per se is largely
234 not concerned with making this easier, although we do have a simple wrapper around `ggplot2` to make
235 it easy to get a quick plot of occurrence data. For example, here we plot 100 occurrences for *Puma*
236 *concolor*.

```
key <- name_backbone(name='Puma concolor')$speciesKey  
dat <- occ_search(taxonKey = key, limit = 100, hasCoordinate = TRUE)  
gbifmap(dat$data)
```



237
238 Another package, `mapr`, is the perfect mapping companion to `rgbif`. It has convenient functions for
239 handling input data from `rgbif`, `spocc`, or arbitrary `data.frame`'s, and output plots for base plots,
240 `ggplot2`, `ggmap` (`ggplot2` with map layers underneath), and interactive maps on GitHub gists or with
241 Leaflet.js.

242 There's no equivalent interface in `pygbif`.

243 GBIF data in other R packages

244 We discuss usage of GBIF data in other R packages throughout the manuscript, but provide a synopsis
245 here for clarity.

246 *taxize*

247 Some of the GBIF taxonomic services are also available in [taxize](#), an R package that focuses on getting
248 data from taxonomic data sources on the web. For example, with `get_gbifid()` one can get GBIF IDs
249 used for a set of taxonomic names - then use those IDs in other functions in **taxize** to get additional
250 information, like taxonomically downstream children.

251 *spocc*

252 GBIF occurrence data is available in the R package [spocc](#) via `rgbif`. **spocc** is a unified interface
253 for fetching species occurrence data from many sources on the web. For example, a user can collect
254 occurrence data from GBIF, iDigBio, and iNaturalist, and easily combine them, then use other packages
255 to clean and visualize the data.

256 **Use cases**

257 The following are three use cases for `rgbif`: niche modeling, spatial change in biodiversity, and
258 distribution mapping.

259 *Ecological niche modeling*

260 In this example, we plot actual occurrence data for *Bradypus* species against a single predictor variable,
261 BIO1 (annual mean temperature). This is only one step in a species distribution modelling workflow.
262 This example can be done using BISON data as well with our `rbison` package.

263 *Load libraries*

```
library("rgbif")  
library("dismo")  
library("maptools")  
library("plyr")
```

264 *Raster files*

265 Make a list of files that are installed with the `dismo` package, then create a `rasterStack` from these

```
files <- list.files(paste(system.file(package = "dismo"), "/ex", sep = ""),
                   "grd", full.names = TRUE)
predictors <- stack(files)
```

266 *Get world boundaries*

```
data(wrld_simpl)
```

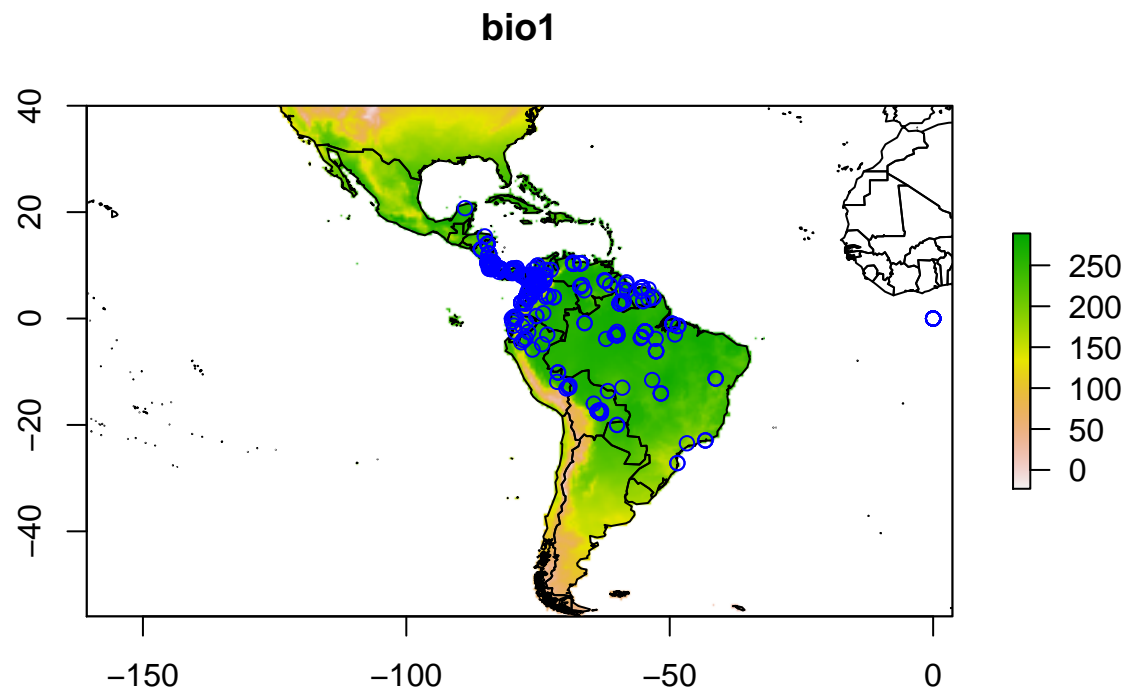
267 *Get GBIF data using the rOpenSci package rgbif*

```
nn <- name_lookup("bradypus*", rank = "species")
nn <- na.omit(unique(nn$data$subKey))
df <- occ_search(taxonKey = nn, hasCoordinate = TRUE, limit = 500)
df <- df[ apply(df, function(x) class(x$data)) %in% "data.frame" ]
df <- ldply(lapply(df, "[", "data"))
df2 <- df[,c('decimalLongitude', 'decimalLatitude')]
```

268 *Plot*

269 (1) Add raster data, (2) Add political boundaries, (3) Add the points (occurrences)

```
plot(predictors, 1)
plot(wrld_simpl, add = TRUE)
points(df2, col = "blue")
```



270

271 *Biodiversity in big cities*

272 In this example, we collect specimen records across different cities using GBIF data from the `rgbif`
273 package.

274 *Load libraries*

```
library("rgbif")  
library("ggplot2")  
library("plyr")  
library("RCurl")  
library("RColorBrewer")
```

275 *Get bounding boxes for some cites*

276 Bounding lat/long data is from [here](#).

```

rawdat <- getURL('https://raw.githubusercontent.com/amyxzhang/boundingbox-cities/master/boundbox
dat <- read.table(text = rawdat, header = FALSE, sep="\t", col.names=c("city","minlat","maxlon",
dat <- data.frame(city=dat$city, minlon=dat$minlon, minlat=dat$minlat, maxlon=dat$maxlon, maxlat=

```

```

getdata <- function(x){
  coords <- as.numeric(x[c('minlon','minlat','maxlon','maxlat')])
  num <- occ_search(geometry = coords)$meta$count
  data.frame(city=x['city'], richness=num, stringsAsFactors = FALSE)
}

```

```

out <- apply(dat, 1, getdata)

```

277 *Merge to original table*

```

out <- merge(dat, ldply(out), by="city")

```

278 *Add centroids from bounding boxes*

```

out <- transform(out, lat = (minlat+maxlat)/2, lon = (minlon+maxlon)/2)

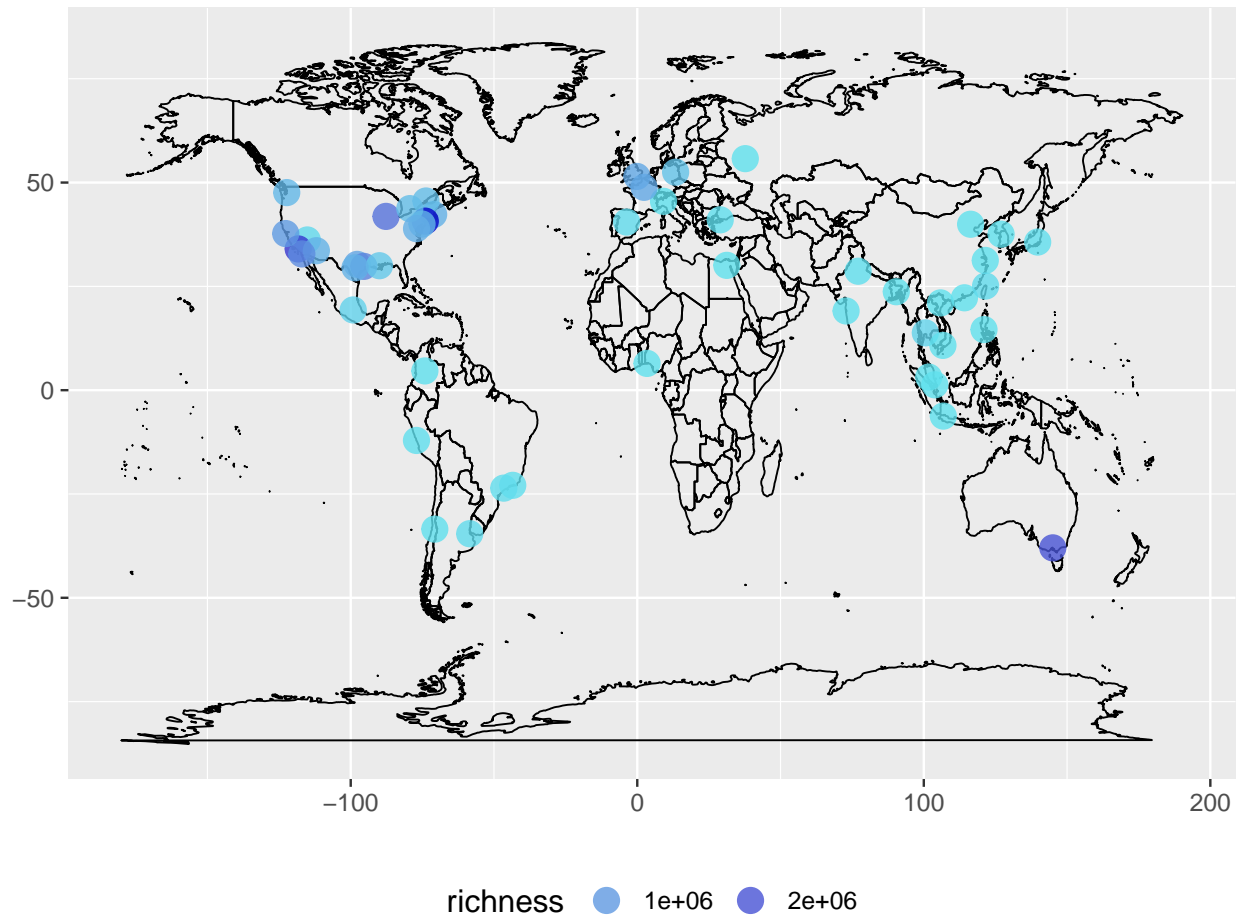
```

279 *Plot data*

```

mapp <- map_data('world')
ggplot(mapp, aes(long, lat)) +
  geom_polygon(aes(group=group), fill="white", alpha=0, color="black", size=0.4) +
  geom_point(data=out, aes(lon, lat, color=richness), size=5, alpha=0.8) +
  scale_color_continuous(low = "#60E1EE", high = "#0404C8") +
  labs(x="", y="") +
  theme_grey(base_size=14) +
  theme(legend.position = "bottom", legend.key = element_blank()) +
  guides(color = guide_legend(keywidth = 2))

```



280

281 *Valley oak occurrence data comparison*

282 This example comes from [Antonio J. Perez-Luque](#) who [shared his plot on Twitter](#). Antonio compared
 283 the occurrences of Valley Oak (*Quercus lobata*) from [GBIF](#) to the distribution of the same species from
 284 the [Atlas of US Trees](#).

285 *Load libraries*

```
library('rgbif')
library('raster')
library('sp')
library('maptools')
library('rgeos')
library('scales')
```

286 *Get GBIF Data for Quercus lobata*

```
keyQ1 <- name_backbone(name='Quercus lobata', kingdom='plants')$speciesKey
dat.Q1 <- occ_search(taxonKey=keyQ1, return='data', limit=50000)
```

287 *Get Distribution map of Q. lobata Atlas of US Trees (Little, E.)*

288 From <http://esp.cr.usgs.gov/data/little/>. And save shapefile in same directory

```
url <- 'http://esp.cr.usgs.gov/data/little/querloba.zip'
tmp <- tempdir()
download.file(url, destfile = "~/querloba.zip")
unzip("~/querloba.zip", exdir = tmp)
ql <- readShapePoly(file.path(tmp, "querloba.shp"))
```

289 *Get Elevation data of US*

```
alt.USA <- getData('alt', country = 'USA')
```

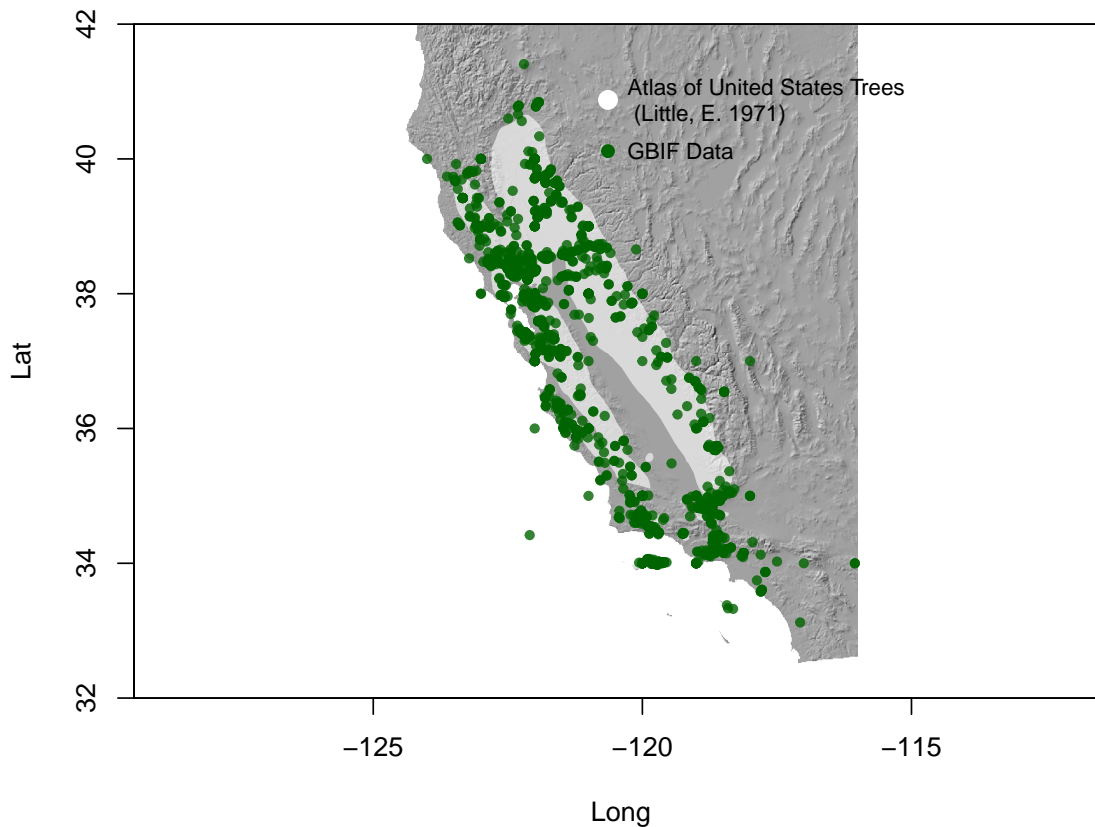
290 *Create Hillshade of US*

```
alt.USA <- alt.USA[[1]]
slope.USA <- terrain(alt.USA, opt = 'slope')
aspect.USA <- terrain(alt.USA, opt = 'aspect')
hill.USA <- hillShade(slope.USA, aspect.USA, angle = 45, direction = 315)
```

291 *Plot map*

```
plot(hill.USA, col = grey(0:100/100), legend = FALSE, xlim = c(-125, -116), ylim = c(32, 42), ma
# add shape from Atlas of US Trees
plot(ql, add = TRUE, col = alpha("white", 0.6), border = FALSE)
# add Gbif presence points
points(dat.Q1$decimalLongitude, dat.Q1$decimalLatitude, cex = .7, pch = 19, col = alpha("darkgreen", 0.5))
legend(x = -121, y = 40.5, "GBIF Data", pch = 19, col = 'darkgreen', bty = 'n', pt.cex = 1, cex = 1.5)
legend(x = -121, y = 41.5, "Atlas of United States Trees \n (Little, E. 1971)", pt.cex = 1.5, cex = 1.5)
```

Distribution of *Quercus lobata*



292

293 Conclusions and future directions

294 The `rgbif` and `pygbif` libraries provide programmatic interfaces to GBIF’s application programming
295 interface (API) - a powerful tool for making research using species occurrence data reproducible. In
296 fact, the `rgbif` package has already been used in 22 scholarly publications (as of 2015-11-14).

297 The `rgbif` package is relatively stable, and should not have many breaking changes unless necessitated
298 due to changes in the GBIF API.

299 The `pygbif` library is still in early development, and will greatly benefit from any feedback and use
300 cases.

301 One area of focus in the future is to attempt to solve many use cases that have been brought up with
302 respect to GBIF data. For example, some specimens are included in GBIF that are located in botanical
303 gardens. For many research questions, researchers are interested in “wild” type occurrences, not those

304 in human curated scenarios. Making removal of these occurrences easy would be very useful, but is
305 actually quite a hard problem.

306 **Acknowledgments**

307 This project was supported in part by the Alfred P Sloan Foundation (Grant No. G-2014-13485), and
308 in part by the Helmsley Foundation (Grant No. 2016PG-BRI004).

309 **Data Accessibility**

310 All scripts and data used in this paper can be found in the permanent data archive Zenodo under
311 the digital object identifier (DOI). This DOI corresponds to a snapshot of the GitHub repository at
312 github.com/sckott/msrgbif. Software can be found at <https://github.com/ropensci/rgbif>, under an
313 MIT license.

314 **References**

- 315 Beck J., Ballesteros-Mejia L., Buchmann CM., Dengler J., Fritz SA., Gruber B., Hof C., Jansen
316 F., Knapp S., Kreft H., Schneider A-K., Winter M., Dormann CF. 2012. Whats on the horizon for
317 macroecology? *Ecography* 35:673–683.
- 318 Brown JH. 1995. *Macroecology*. University of Chicago Press.
- 319 Brown KA., Parks KE., Bethell CA., Johnson SE., Mulligan M. 2015. Predicting plant diversity patterns
320 in madagascar: Understanding the effects of climate and land cover change in a biodiversity hotspot.
321 *PLOS ONE* 10:e0122721.
- 322 Ceballos G., Ehrlich PR., Barnosky AD., Garcia A., Pringle RM., Palmer TM. 2015. Accelerated
323 modern human-induced species losses: Entering the sixth mass extinction. *Science Advances* 1:e1400253–
324 e1400253.
- 325 Chamberlain S., Ram K., Barve V., Mcglinn D. *Pygbif: Interface to the global 'biodiversity' information*
326 *facility 'aPI'*.
- 327 Chamberlain S. *Rgbif: Python client for gBIF*.

328 Faulkner KT., Robertson MP., Rouget M., Wilson JR. 2014. A simple, rapid methodology for developing
329 invasive species watch lists. *Biological Conservation* 179:25–32.

330 Febbraro MD., Lurz PWW., Genovesi P., Maiorano L., Girardello M., Bertolino S. 2013. The use of
331 climatic niches in screening procedures for introduced species to evaluate risk of spread: A case with
332 the american eastern grey squirrel. *PLoS ONE* 8:e66559.

333 Ferretti F., Verd GM., Seret B., Šprem JS., Micheli F. 2015. Falling through the cracks: The fading
334 history of a large iconic predator. *Fish and Fisheries*:n/a–n/a.

335 Ficetola GF., Rondinini C., Bonardi A., Baisero D., Padoa-Schioppa E. 2014. Habitat availability for
336 amphibians and extinction threat: A global analysis. *Diversity and Distributions* 21:302–311.

337 Herring J. 2011. OpenGIS implementation standard for geographic information-simple feature access-
338 part 1: Common architecture. *OGC Document* 4:122–127.

339 María Mendoza., Ospina OE., Cárdenas-Henao H., García-R JC. 2015. A likelihood inference of
340 historical biogeography in the world’s most diverse terrestrial vertebrate genus: Diversification of
341 direct-developing frogs (craugastoridae: Pristimantis) across the neotropics. *Molecular Phylogenetics*
342 *and Evolution* 85:50–58.

343 Pimm SL., Jenkins CN., Abell R., Brooks TM., Gittleman JL., Joppa LN., Raven PH., Roberts CM.,
344 Sexton JO. 2014. The biodiversity of species and their rates of extinction, distribution, and protection.
345 *Science* 344:1246752–1246752.

346 R Core Team. 2014. *R: A language and environment for statistical computing*. Vienna, Austria: R
347 Foundation for Statistical Computing.